

# **CSE 31**

# **Computer Organization**

Lecture 9 – C Memory Management (cont.)

# Announcements

- Labs
  - Lab 3 due this week (**with 7 days grace period** after due date)
    - » Demo is REQUIRED to receive full credit
  - Lab 4 out this week
    - » Due at 11:59pm on the same day of your next lab (with 7 days grace period after due date)
    - » You must demo your submission to your TA within 14 days from posting of lab
    - » Demo is REQUIRED to receive full credit
- Reading assignments
  - Reading 02 (zyBooks 2.1 – 2.9) due 22-FEB
    - » Complete **Participation Activities** in each section to receive grade towards Participation
    - » IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses
- Homework assignment
  - Homework 01 (zyBooks 1.1 – 1.5) due 22-FEB
    - » Complete **Challenge Activities** in each section to receive grade towards Homework
    - » IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses

# Announcements

- Project 01
  - Due 17-MAR
  - Can work in teams of 2 students
    - » Each team member must identify teammate in “Comments...” text-box at the submission page
    - » If working in teams, each student must submit code (can be the same as teammate) and demo individually
    - » Grade can vary among teammates depending on demo
  - Demo required for project grade
    - » No partial credit for submission without demo
  - No grace period
    - » Must complete submission and demo by due date.

# The Heap (Dynamic memory)

- Large pool of memory, not allocated in contiguous order
  - back-to-back requests for heap memory could result in blocks very far apart
  - where Java/C++ **new** command allocates memory
- In C, specify number of bytes of memory explicitly to allocate item

```
int *ptr;  
ptr = (int *) malloc(10*sizeof(int));  
/* malloc returns type (void *),  
so need to cast to right type */
```

– `malloc()`: Allocates raw, uninitialized memory from heap

# Memory Management

- How do we manage memory?
  - Code/Text, Static
    - » Simple
    - » They never grow or shrink
  - Stack
    - » Simple
    - » Stack frames are created and destroyed in last-in, first-out (LIFO) order
  - Heap
    - » Tricky
    - » Memory can be allocated / deallocated at any time


# Heap Management Requirements

- Want `malloc()` and `free()` to run quickly.
- Want minimal memory overhead
- Want to avoid fragmentation\*
  - When most of our free memory is in many small chunks
  - In this case, we might have many free bytes but not be able to satisfy a large request since the free bytes are not contiguous in memory.

\* This is technically called *external fragmentation*


# Heap Management

- An example
  - Request R1 for 100 bytes
  - Request R2 for 1 byte
  - Memory from R1 is freed



# Heap Management

- An example
  - Request R1 for 100 bytes
  - Request R2 for 1 byte
  - Memory from R1 is freed
  - Request R3 for 50 bytes





# K&R Malloc/Free Implementation

- From Section 8.7 of K&R
  - Code in the book uses some C language features we haven't discussed and is written in a very terse style, don't worry if you can't decipher the code
- Each block of memory is preceded by a header that has two fields:
  - size of the block
  - a pointer to the next block
- All free blocks are kept in a circular linked list, the pointer field is unused in an allocated block

# K&R Implementation

- `malloc()` searches the free list for a block that is big enough. If none is found, more memory is requested from the operating system. If what it gets can't satisfy the request, it fails.
- `free()` checks if the blocks adjacent to the freed block are also free
  - If so, adjacent free blocks are merged (**coalesced**) into a single, larger free block
  - Otherwise, the freed block is just added to the free list

# Choosing a block in `malloc()`

- If there are multiple free blocks of memory that are big enough for some request, how do we choose which one to use?
  - **best-fit**: choose the smallest block that is big enough for the request
  - **first-fit**: choose the first block we see that is big enough
  - **next-fit**: like first-fit but remember where we finished searching and resume searching from there

# Tradeoffs of allocation policies

- **Best-fit:** Tries to limit fragmentation but at the cost of time (must examine all free blocks for each malloc).
  - Leaves lots of small blocks (why?)
- **First-fit:** Quicker than best-fit (why?) but potentially more fragmentation.
  - Tends to concentrate small blocks at the beginning of the free list (why?)
- **Next-fit:** Does not concentrate small blocks at front like first-fit, should be faster as a result.

## Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

	<b>1</b>	<b>2</b>	<b>3</b>
a)	F	F	T
b)	F	T	T
c)	T	F	F
d)	T	F	T
e)	T	T	T

## Quiz – Pros and Cons of fits

- 1) **first-fit** results in many **small blocks** at the beginning of the free list
- 2) **next-fit** is **slower than first-fit**, since it takes longer in steady state to find a match
- 3) **best-fit** leaves lots of tiny blocks

	<b>1</b>	<b>2</b>	<b>3</b>
a)	F	F	T
b)	F	T	T
c)	T	F	F
<b>d)</b>	<b>T</b>	<b>F</b>	<b>T</b>
e)	T	T	T

# Summary

- C has 3 pools of memory
  - Static storage: global variable storage, basically permanent, entire program run
  - The Stack: local variable storage, parameters, return address
  - The Heap (dynamic storage): `malloc()` grabs space from here, `free()` returns it.
- `malloc()` handles free space with freelist. Three different ways to find free space when given a request:
  - **First fit** (find first one that's free)
  - **Next fit** (same as first, but remembers where left off)
  - **Best fit** (finds most “snug” free space)

# Slab Allocator


- A different approach to memory management (used in GNU `libc`)
- Divide blocks in to “large” and “small” by picking a threshold size (say 128kB). Blocks larger than this threshold are managed with a **freelist** (as before).
- For small blocks, allocate blocks in sizes that are powers of 2
  - e.g., if program wants to allocate 20 bytes, actually give it 32 bytes



# Slab Allocator

- Bookkeeping for small blocks is relatively easy
  - Use a **bitmap** for each range of blocks of the same size
- Allocating is easy and fast
  - Compute the size of the block to allocate and find a free bit in the corresponding bitmap.
- Freeing is also easy and fast
  - Figure out which slab the address belongs to and clear the corresponding bit.

# Slab Allocator



16 byte block bitmap: 11011000

32 byte block bitmap: 0111

64 byte block bitmap: 00

# Slab Allocator Tradeoffs

- Extremely fast for small blocks.
- Slower for large blocks
  - But presumably the program will take more time to do something with a large block, so the overhead is not as critical.
- Minimal space overhead
- No external fragmentation (as we defined it before)
  - For small blocks, but still have wasted space!

# Internal vs. External Fragmentation


- With the slab allocator, difference between requested size and next power of 2 is wasted
  - e.g., if program wants to allocate 20 bytes and we give it a 32 byte block, 12 bytes are unused.
- We also refer to this as fragmentation but call it **internal fragmentation** since the wasted space is actually within an allocated block
- **External fragmentation**: wasted space between allocated blocks

# Buddy System

- Yet another memory management technique (used in Linux kernel)
- Like GNU's "slab allocator", but only allocate blocks in sizes that are powers of 2 (internal fragmentation is possible)
- Initially entire memory space treated as single block whose size is power of 2
  - If **request size** > **0.5 \* initial block size**, allocate entire block
  - Otherwise, split block in 2 companion buddies
  - If **request size** > **0.5 \* buddy size**, allocate entire buddy
  - Otherwise, split one buddy in half again
  - Repeat till smallest block  $\geq$  size of request is found

# Buddy System

- If no free block of size  $n$  is available, find a block of size  $2n$  and split it into two blocks of size  $n$
- When a block of size  $n$  is freed, if its neighbor of size  $n$  is also free, combine the blocks into a single block of size  $2n$ 
  - **Buddy** is a block in the other half of a larger block



- Same speed advantages as slab allocator

# Buddy System

Buddy system at work, considering a 1024k (1-megabyte) initial block and the process requests as shown at the left of the table.

	0	128k	256k	512k	1024k
start	1024k				
A=70K	A	128	256	512	
B=35K	A	B 64	256	512	
C=80K	A	B 64	C 128	512	
A ends	128	B 64	C 128	512	
D=60K	128	B D	C 128	512	
B ends	128	64 D	C 128	512	
D ends	256		C 128	512	
C ends	512			512	
end	1024k				

# Allocation Schemes

- So which memory management scheme (K&R, slab, buddy) is best?
  - There is no single best approach for every application.
  - Different applications have different allocation / deallocation patterns.
  - A scheme that works well for one application may work poorly for another application.



# Automatic Memory Management

- Dynamically allocated memory is difficult to track
  - Why not track it **automatically**?
- If we can keep track of what memory is in use, we can reclaim everything else.
  - Unreachable memory is called **garbage**, the process of reclaiming it is called **garbage collection**.
- So how do we track what is in use?

# Tracking Memory Usage

- Techniques depend heavily on the programming language and rely on help from the compiler.
- Start with all pointers in global variables and local variables ([root set](#)).
- Recursively examine dynamically allocated objects we see a pointer to.
  - We can do this in **constant space** by reversing the pointers on the way down
- We will cover 3 schemes to collect garbage