

CSE 31

Computer Organization

Lecture 19 – MIPS Logical Operators,
Multiplication and Division, Instruction Format

Announcements

- Labs

- Lab 6 grace period ends this week
 - » No penalty for submission during grace period
 - » Demo is REQUIRED to receive full credit
- Lab 7 due this week (with **7 days grace period** after due date)
 - » Demo is REQUIRED to receive full credit
- Lab 8 out this week
 - » Due at 11:59pm on the same day of your lab after next (with 7 days grace period after due date)
 - » You must demo your submission to your TA within 14 days from posting of lab
 - » Demo is REQUIRED to receive full credit

- Reading assignments

- Reading 06 (zyBooks 6.4 – 6.7) due 10-APR
 - » Complete **Participation Activities** in each section to receive grade
 - » IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses

Announcements

- Homework assignment
 - Homework 05 (zyBooks 1.6 – 1.7) due 10-APR
 - » Complete **Challenge Activities** in each section to receive grade
 - » IMPORTANT: Make sure to submit score to CatCourses by using the link provided on CatCourses
- **Midterm 02 postponed to 19-APR**
- Project 02
 - Due 05-MAY
 - Can work in teams of 2 students
 - » Each team member must identify teammate in “Comments...” text-box at the submission page
 - » If working in teams, each student must submit code (can be the same as teammate) and demo individually
 - » Grade can vary among teammates depending on demo
 - Demo required for project grade
 - » No partial credit for submission without demo
 - **No grace period**
 - » **Must complete submission and demo by due date.**

Bitwise Operations

- So far, we've done arithmetic (**add**, **sub**, **addi**), mem access (**lw** and **sw**), & branches and jumps.
- All of these instructions view contents of register as a single quantity (e.g., signed or unsigned int)
- **New Perspective**: View register as 32 raw bits rather than as a single 32-bit number
 - Since registers are composed of 32 bits, wish to access individual bits (or groups of bits) rather than the whole.
- Introduce two new classes of instructions
 - Logical & Shift Ops

Logical Operators (1/3)

- Two basic logical operators:
 - AND: outputs 1 only if **all** inputs are 1
 - OR: outputs 1 if **at least one** input is 1
- Truth Table: standard table listing all possible combinations of inputs and resultant output

A	B	A AND B	A OR B
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	1

Logical Operators (2/3)

- Logical Instruction Syntax:

Format: 1 2, 3, 4

where:

- 1) operation name
 - 2) register that will receive value
 - 3) first operand (register)
 - 4) second operand (register) or immediate (numerical constant)
- In general, can define them to accept > 2 inputs, but in the case of MIPS assembly, these accept exactly 2 inputs and produce 1 output
 - Again, rigid syntax, simpler hardware

Logical Operators (3/3)

- Instruction Names:
 - **and, or**: Both of these expect the third argument to be a register
 - **andi, ori**: Both of these expect the third argument to be an immediate
- MIPS Logical Operators are all **bitwise**, meaning that bit 0 of the output is produced by the respective bit 0's of the inputs, bit 1 by the bit 1's, etc.
 - C: Bitwise AND is **&** (e.g., **z = x & y;**)
 - C: Bitwise OR is **|** (e.g., **z = x | y;**)

Uses of Logical Operators (1/3)

- Note that **anding** a bit with 0 produces a 0 at the output while **anding** a bit with 1 produces the original bit.
- This can be used to create a **mask**.

– Example:

1011 0110 1010 0100 0011

0000 0000 0000 0000 0000

– The result of **anding** these:

0000 0000 0000 0000 0000

mask:

1101 1001 1010

1111 1111 1111

1101 1001 1010

mask last 12 bits

Uses of Logical Operators (2/3)

- The second bitstring in the example is called a **mask**. It is used to isolate the rightmost 12 bits of the first bitstring by masking out the rest of the string (e.g. setting to all 0s).
- Thus, the **and** operator can be used to set certain portions of a bitstring to 0s, while leaving the rest alone.
 - In particular, if the first bitstring in the above example were in `$t0`, then the following instruction would mask it:
`andi $t0, $t0, 0xFFF`

Uses of Logical Operators (3/3)

- Similarly, note that **oring** a bit with **1** sets a **1** at the output while **oring** a bit with **0** keeps the original bit.
- Often used to force certain bits to **1s**.
 - For example, if `$t0` contains `0x12345678`, then after this instruction:
`ori $t0, $t0, 0xFFFF`
... `$t0` will contain `0x1234FFFF`
» (i.e., the high-order 16 bits are untouched, while the low-order 16 bits are forced to **1s**).

Shift Instructions (review) (1/4)

- Move (shift) all the bits in a word to the left or right by a number of bits.

– Example: shift right by 8 bits

1001 0010 0011 0100 0101 0110 0111 1000

0000 0000 1001 0010 0011 0100 0101 0110

– Example: shift left by 8 bits

0001 0010 1011 0100 0101 0110 0111 1000

1011 0100 0101 0110 0111 1000 0000 0000

Shift Instructions (2/4)

- Shift Instruction Syntax:

Format: 1 2, 3, 4

where:


- 1) operation name
- 2) register that will receive value
- 3) first operand (register)
- 4) shift amount (constant < 32)

- MIPS shift instructions:


1. **sll** (shift left logical): shifts left and fills emptied bits with 0s
2. **srl** (shift right logical): shifts right and fills emptied bits with 0s
3. **sra** (shift right arithmetic): shifts right and fills emptied bits by sign extending

Shift Instructions (3/4)

- Example: shift right arithmetic (sra) by 8 bits



- Example: shift right arithmetic (sra) by 8 bits



Shift Instructions (4/4)

- Since shifting is faster than multiplication, a good compiler usually notices when C code multiplies by a power of 2 and compiles it to a shift instruction:

`a *= 8;` (in C)

would compile to:

`sll $s0, $s0, 3` (in MIPS)

- Likewise, shift right to divide by powers of 2 (rounds towards $-\infty$)
 - remember to use `sra`

Summary

- Logical and Shift Instructions
 - Operate on bits individually, unlike arithmetic, which operate on entire word.
 - Use to isolate fields, either by masking or by shifting back and forth.
 - Use shift left logical, **sll** for multiplication by powers of 2
 - Use shift right logical, **srl** for division by powers of 2 of unsigned numbers (**unsigned int**)
 - Use shift right arithmetic, **sra** for division by powers of 2 of signed numbers (**int**)
- New Instructions:
and, andi, or, ori, sll, srl, sra

Integer Multiplication (1/3)

- Paper and pencil example (unsigned):

Multiplicand	1000	8
Multiplier	$\times \begin{array}{c} \text{blue}1\text{orange}0\text{green}0\text{red}1 \end{array}$	9
	\hline	
	1000	
	0000	
	0000	
	+1000	
	\hline	
	01001000	

- m bits \times n bits = $m + n$ bit product

Integer Multiplication (2/3)

- In MIPS, we multiply registers, so:
 - 32-bit value x 32-bit value = 64-bit value
- Syntax of Multiplication (signed):
 - `mult` register1, register2 **No destination register!**
 - Multiplies 32-bit values in those registers & puts 64-bit product in special result regs:
 - » puts product **upper half in hi**, **lower half in lo**
 - **hi** and **lo** are 2 registers separate from the 32 general purpose registers
 - Use **mfhi register** and **mflo register** to move from **hi**, **lo** to another register

Integer Multiplication (3/3)

- Example:

- in C: `a = b * c;`

- in MIPS:

- » let `b` be `$s2`; let `c` be `$s3`; and let `a` be `$s0` and `$s1` (since it may be up to 64 bits)

- `mult $s2, $s3 # b*c`

- `mfhi $s0 # upper half of product into $s0`

- `mflo $s1 # lower half of product into $s1`

- Note: Often, we only care about the lower half of the product.

Integer Division (1/2)

- Paper and pencil example (unsigned):

		<u>1001</u>	Quotient
Divisor	1000	1001010	Dividend
		- 1000	
		10	
		101	
		1010	
		- 1000	
		10	Remainder
			(or Modulo result)

- $\text{Dividend} = \text{Quotient} \times \text{Divisor} + \text{Remainder}$

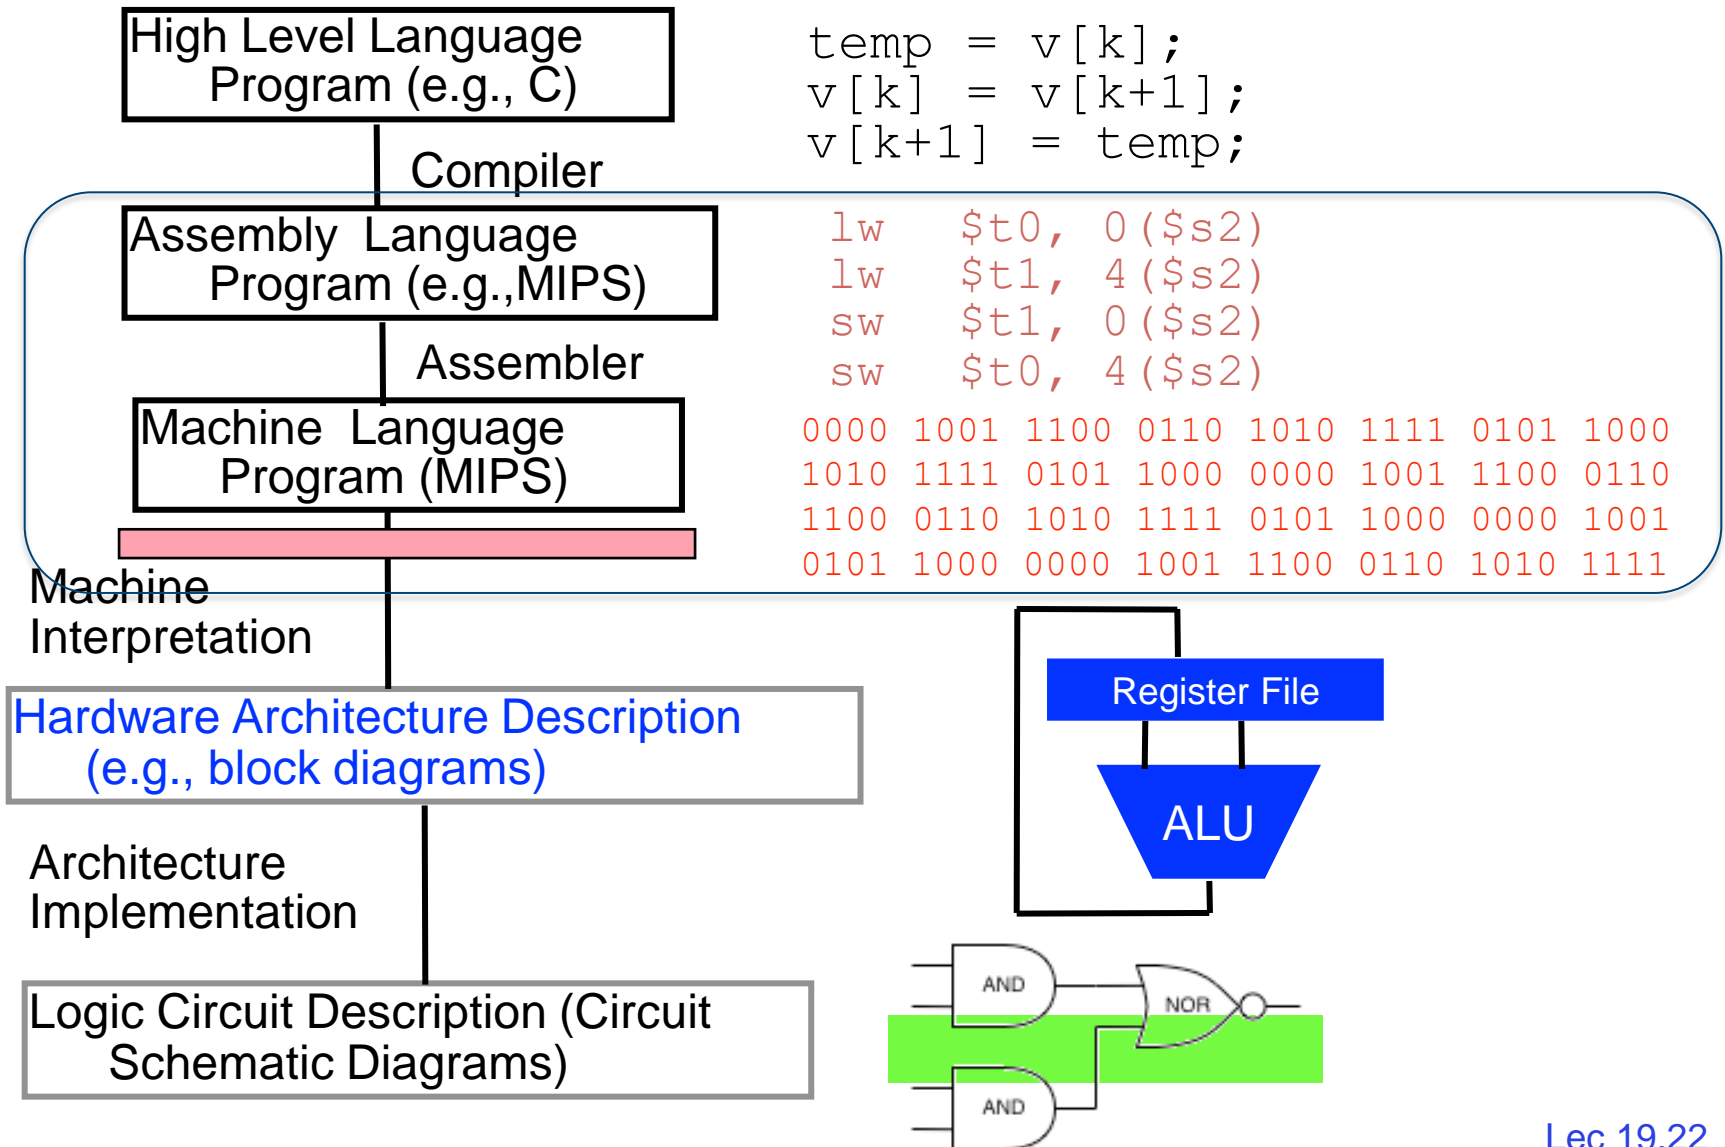
Integer Division (2/2)

- Syntax of Division (signed):
 - `div register1, register2`
 - Divides 32-bit register1 by 32-bit register2
 - Puts remainder of division in hi, quotient in lo
- Implements C division (/) and modulo (%)
- Example in C: `a = c / d; b = c % d;`
- in MIPS: `a=$s0; b=$s1; c=$s2; d=$s3`

<code>div</code>	<code>\$s2, \$s3</code>	<code># lo=c/d, hi=c%d</code>
<code>mflo</code>	<code>\$s0</code>	<code># get quotient</code>
<code>mfhi</code>	<code>\$s1</code>	<code># get remainder</code>

Machine Language

Levels of Representation (abstractions)



Big Idea: Stored-Program Concept

- Where are programs stored when they are being run?
 - How are they stored in memory?
- Computers built on 2 key principles:
 - Instructions are represented as bit patterns - can think of these as numbers.
 - Therefore, entire programs can be stored in memory to be read or written just like data.
- Simplifies SW/HW of computer systems:
 - Memory technology for data also used for programs

Consequence #1: Everything Addressed

- Since all instructions and data are stored in memory, everything has a memory address: instructions, data words
 - both branches and jumps use these
- C pointers are just memory addresses: they can point to anything in memory
 - Unconstrained use of addresses can lead to nasty bugs; up to you in C; limits in Java
- One register keeps address of instruction being executed: “Program Counter” (PC)
 - Basically, a pointer to memory: Intel calls it Instruction Address Pointer, a better name

Consequence #2: Binary Compatibility

- Programs are distributed in binary form
 - Programs bound to specific instruction set
 - Different version for **Macs** and **PCs**
- New machines want to run old programs (“binaries”) as well as programs compiled to new instructions
 - Leads to “backward compatible” instruction set evolving over time
 - Selection of Intel 8086 in 1981 for 1st IBM PC is major reason latest PCs still use 80x86 instruction set; could still run program from 1981 PC today

Instructions as Numbers (1/2)

- Currently all data we work with is in words (32-bit blocks):
 - Each register is a word.
 - **lw** and **sw** both access memory one word at a time.
- So how do we represent instructions?
 - Remember: Computer only understands 1s and 0s, so “**add \$t0, \$0, \$0**” is meaningless.
 - MIPS wants simplicity: since data is in words, make instructions into words too!

Instructions as Numbers (2/2)

- One word is 32 bits, so divide instruction word into “fields”.
- Each field tells processor something about the instruction.
- We could define different fields for each instruction, but MIPS is based on simplicity, so define 3 basic types of instruction formats:
 - I-format
 - J-format
 - R-format
 - What do these letters (I, J, R) stand for?

Instruction Formats

- **3 formats**
- **I-format**: used for instructions with immediates, **lw** and **sw** (since offset counts as an immediate), and branches (**beq** and **bne**),
 - (but not the shift instructions; later)
- **J-format**: used for **j** and **jal**
- **R-format**: used for all other instructions
- **Why 3 different formats?**
 - It will soon become clear why the instructions have been partitioned in this way.

R-Format Instructions (1/5)

- Define “fields” of the following number of bits each: $6 + 5 + 5 + 5 + 5 + 6 = 32$

6	5	5	5	5	6
---	---	---	---	---	---

- For simplicity, each field has a name:

opcode	rs	rt	rd	shamt	funct
--------	----	----	----	-------	-------

- Important: On these slides and in book, each field is viewed as a 5- or 6-bit unsigned integer, not as part of a 32-bit integer.
 - Consequence: 5-bit fields can represent any number 0-31, while 6-bit fields can represent any number 0-63.

R-Format Instructions (2/5)

- What do these field integer values tell us?
 - **opcode**: partially specifies what instruction it is
 - » **Note: This number is equal to 0 for all R-Format instructions.**
 - **funct**: combined with `opcode`, this number exactly specifies the instruction
- Question: Why aren't `opcode` and `funct` a single 12-bit field?
 - We'll answer this later.

R-Format Instructions (3/5)

- More fields:
 - rs (Source Register): **generally** used to specify register containing **first operand**
 - rt (Target Register): **generally** used to specify register containing **second operand** (note that name is misleading)
 - rd (Destination Register): **generally** used to specify register which will **receive result** of computation

R-Format Instructions (4/5)

- Notes about register fields:
 - Do we need more than 5 bits?
 - Each register field is exactly 5 bits, which means that it can specify any unsigned integer in the range 0-31. Each of these fields specifies one of the 32 registers by number.
 - The word “generally” was used because there are exceptions
 - » **mult** and **div** have nothing important in the **rd** field since the dest registers are **hi** and **lo**
 - » **mfhi** and **mflo** have nothing important in the **rs** and **rt** fields since the source is determined by the instruction.

R-Format Instructions (5/5)

- Final field:
 - shamt: This field contains the amount a shift instruction will shift by. Shifting a 32-bit word by more than 31 is useless, so this field is only 5 bits (so it can represent the numbers 0-31).
 - **This field is set to 0 in all but the shift instructions.**
- For a detailed description of field usage for each instruction, see MIPS_Reference_Sheet.pdf in CatCourses (You will be provided in all exams)